



Recap

2

- ▶ Data processing trend
- ▶ Big Data and its characteristics
- ▶ Applications of big data
- ▶ Various computing technologies
- ▶ History of Hadoop
- ▶ RDBMS vs Hadoop
- ▶ Major components of Hadoop cluster
- ▶ Unix commands

Agenda for today

3

- ▶ The Hadoop Distributed File System
- ▶ Accessing HDFS through CLI
- ▶ MapReduce detailed discussion
- ▶ Running your first MapReduce program

HDFS

- ▶ HDFS is a file system designed for storing very large files with streaming data access patterns, running on clusters of commodity hardware

Large Files

**Streaming data
access**

**Commodity
hardware**

HDFS Blocks

- ▶ Single unit of storage
- ▶ Default block size is 128MB
- ▶ Size of block will drive the ratio of time to read a block to the seek for a block

Exercise

6

- ▶ What should be the block size to make seek time 1% of read time for given hardware configuration
 - Seek time: 10 ms
 - Data read rate: 100MB/s
- ▶ Solution:
 - Let say x MB is the block size then read time = $x/100$ seconds
 - To fulfill the given condition 1% of $x/100 = 10$ ms
 - Hence $x=100$ MB


Benefits of blocks

7

- ▶ Files can be larger than a single disk
- ▶ Simplicity at storage level as data node doesn't store any metadata
- ▶ Fault tolerance by replicating blocks

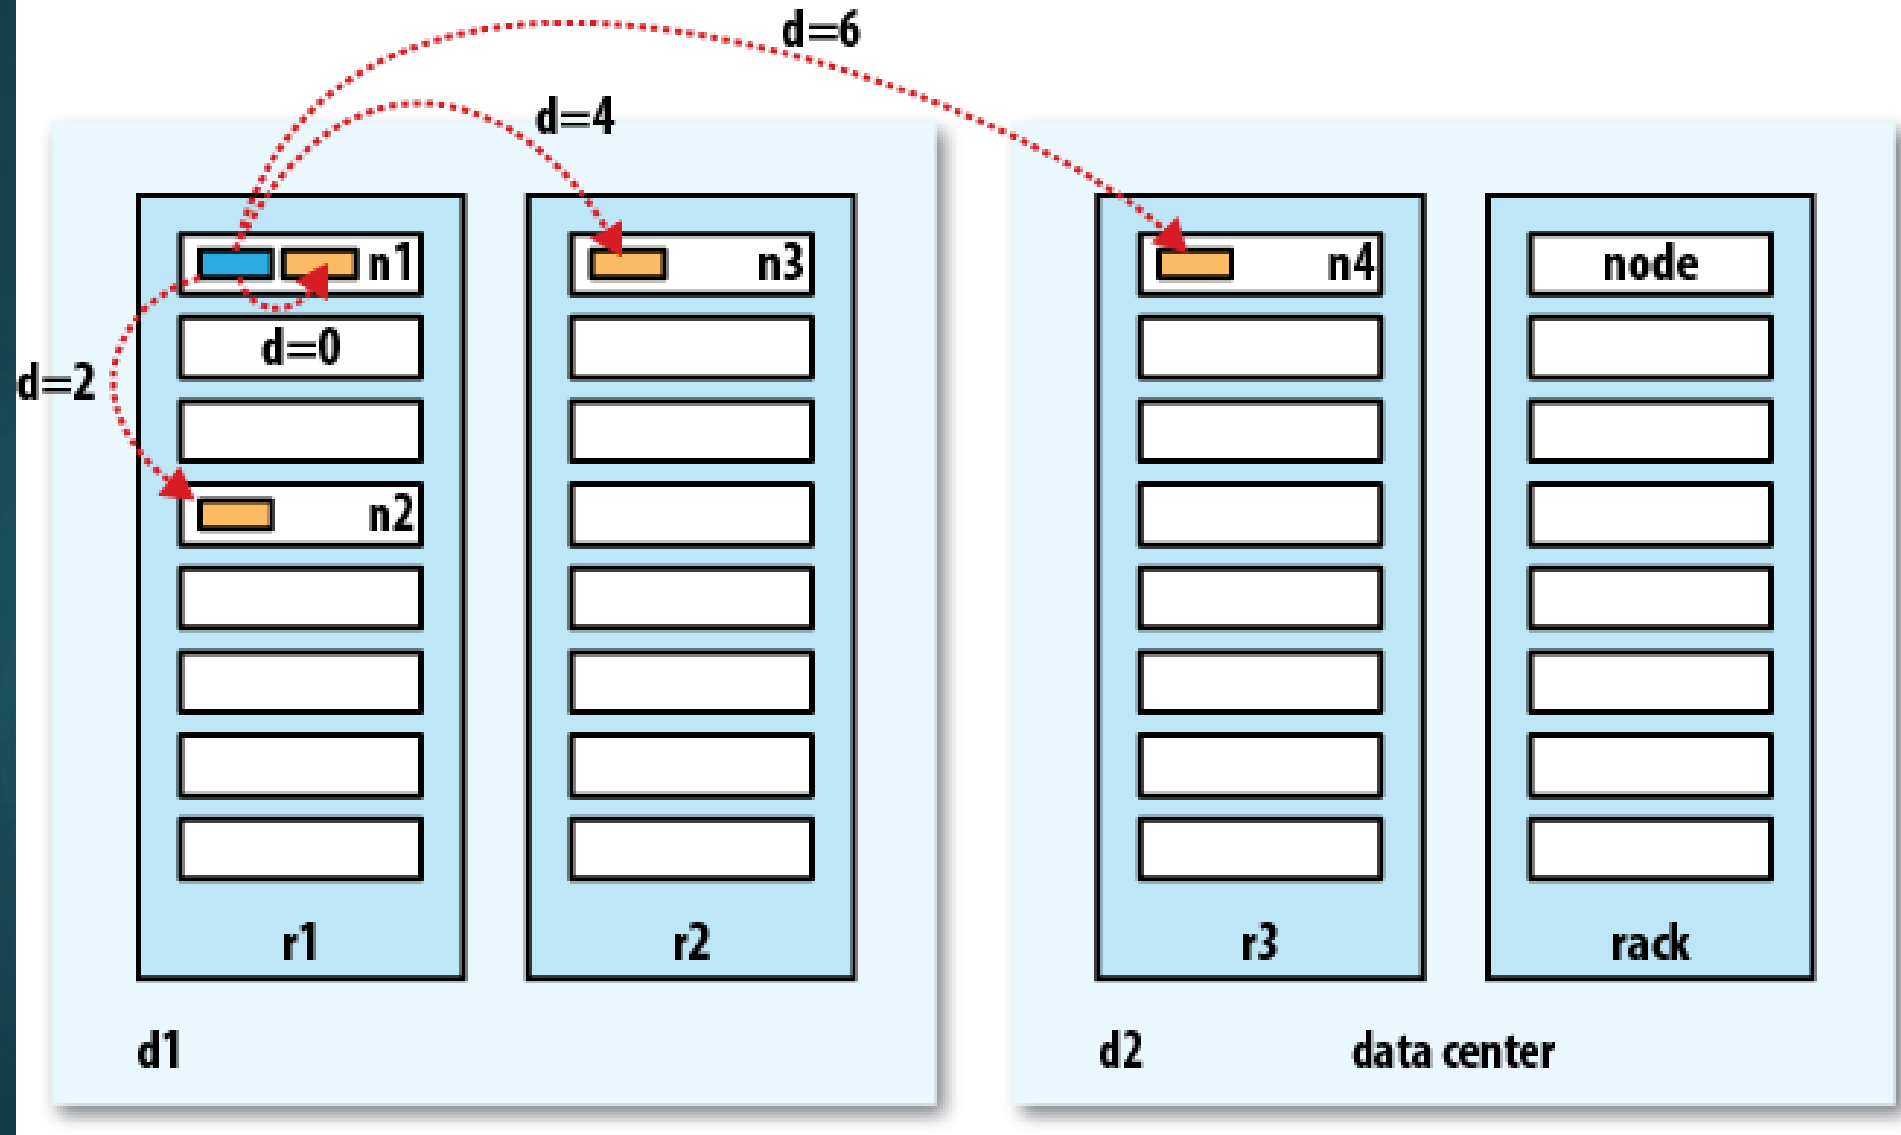
File system metadata


- ▶ Namenode stores the metadata
- ▶ Backup of metadata on secondary namenode



Network Topology

9





data center

Rack
awareness

HDFS CLI Copy-To-From command

11

```
hadoop fs -copyFromLocal <source> <target>
```

- ▶ To copy myTextFile.txt from local unix file system to HDFS under data directory

```
hadoop fs -copyFromLocal myTextFile.txt /data/
```

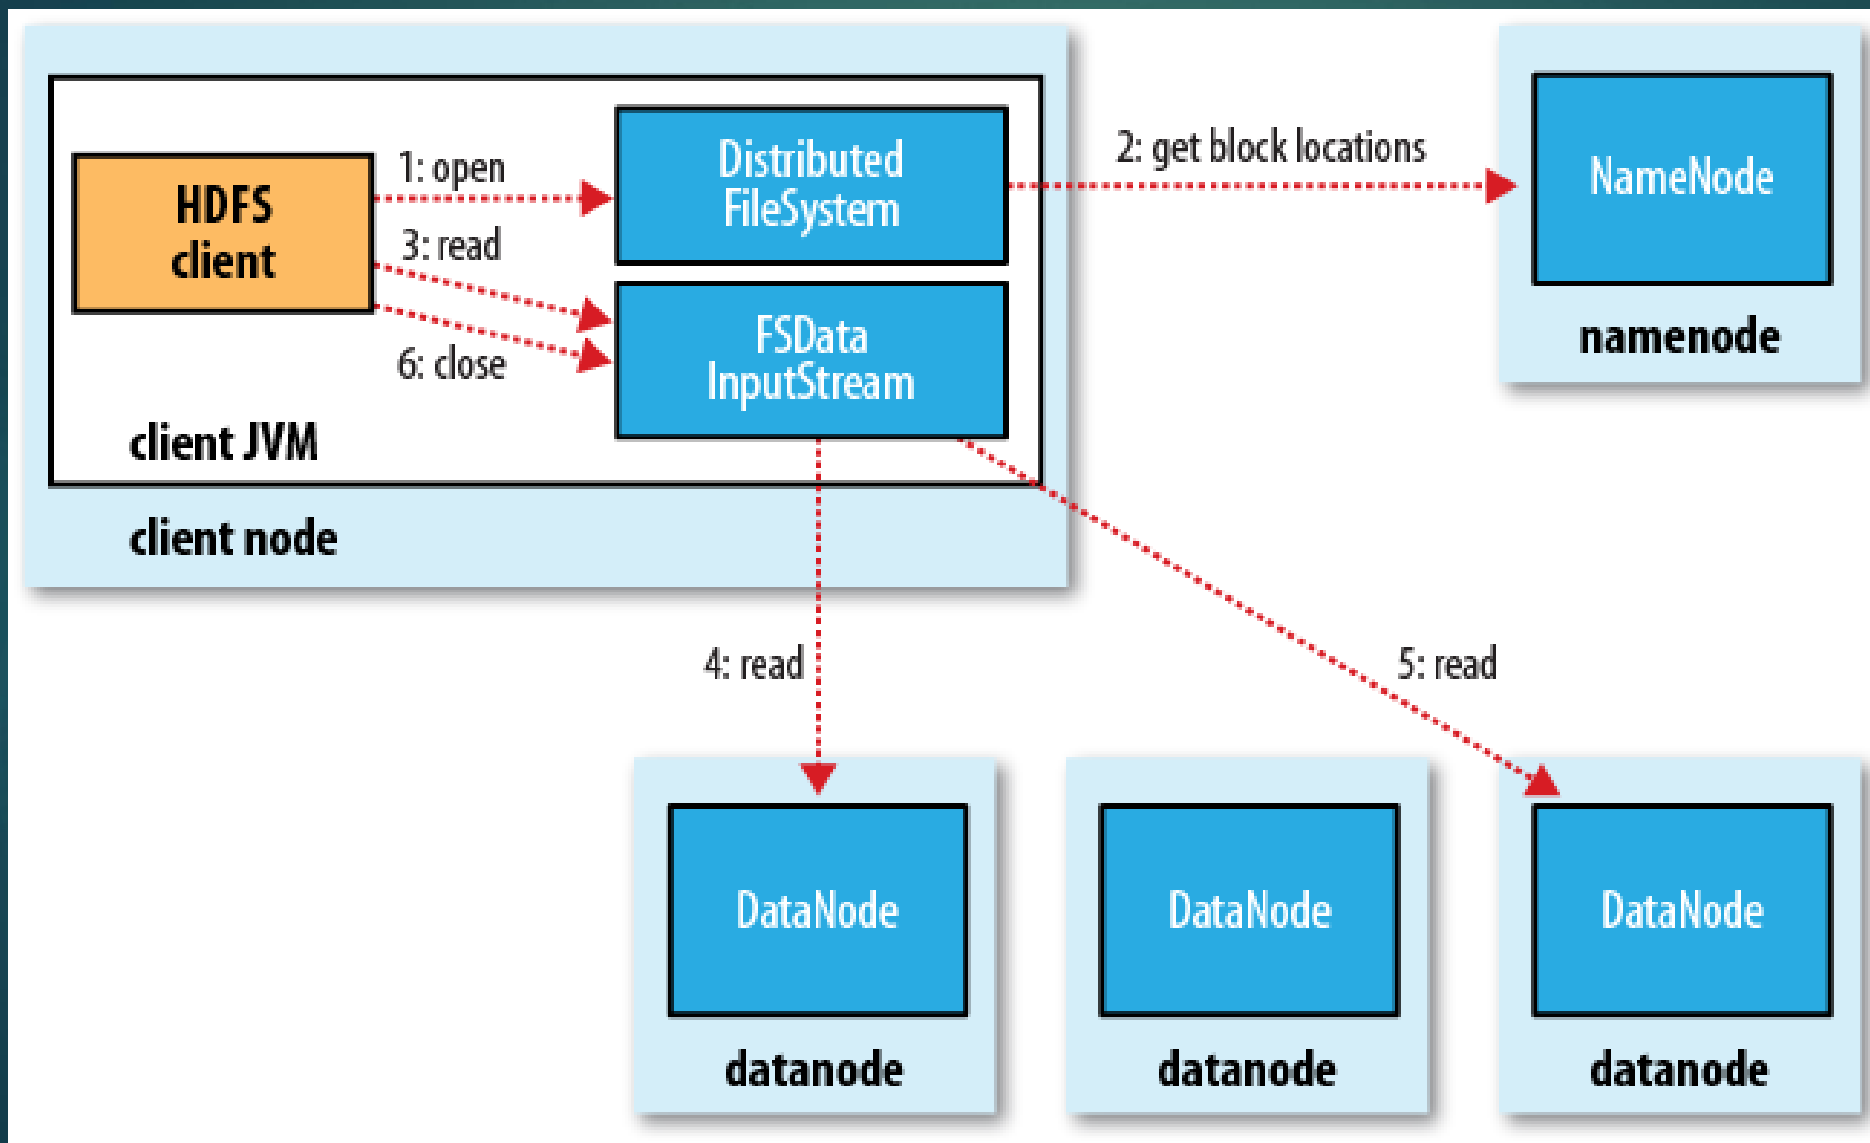
```
hadoop fs -copyToLocal <source> <target>
```

- ▶ To copy myTextFile.txt from data directory in HDFS to local file system on Desktop

```
hadoop fs -copyToLocal /data/myTextFile.txt ~/Desktop/
```

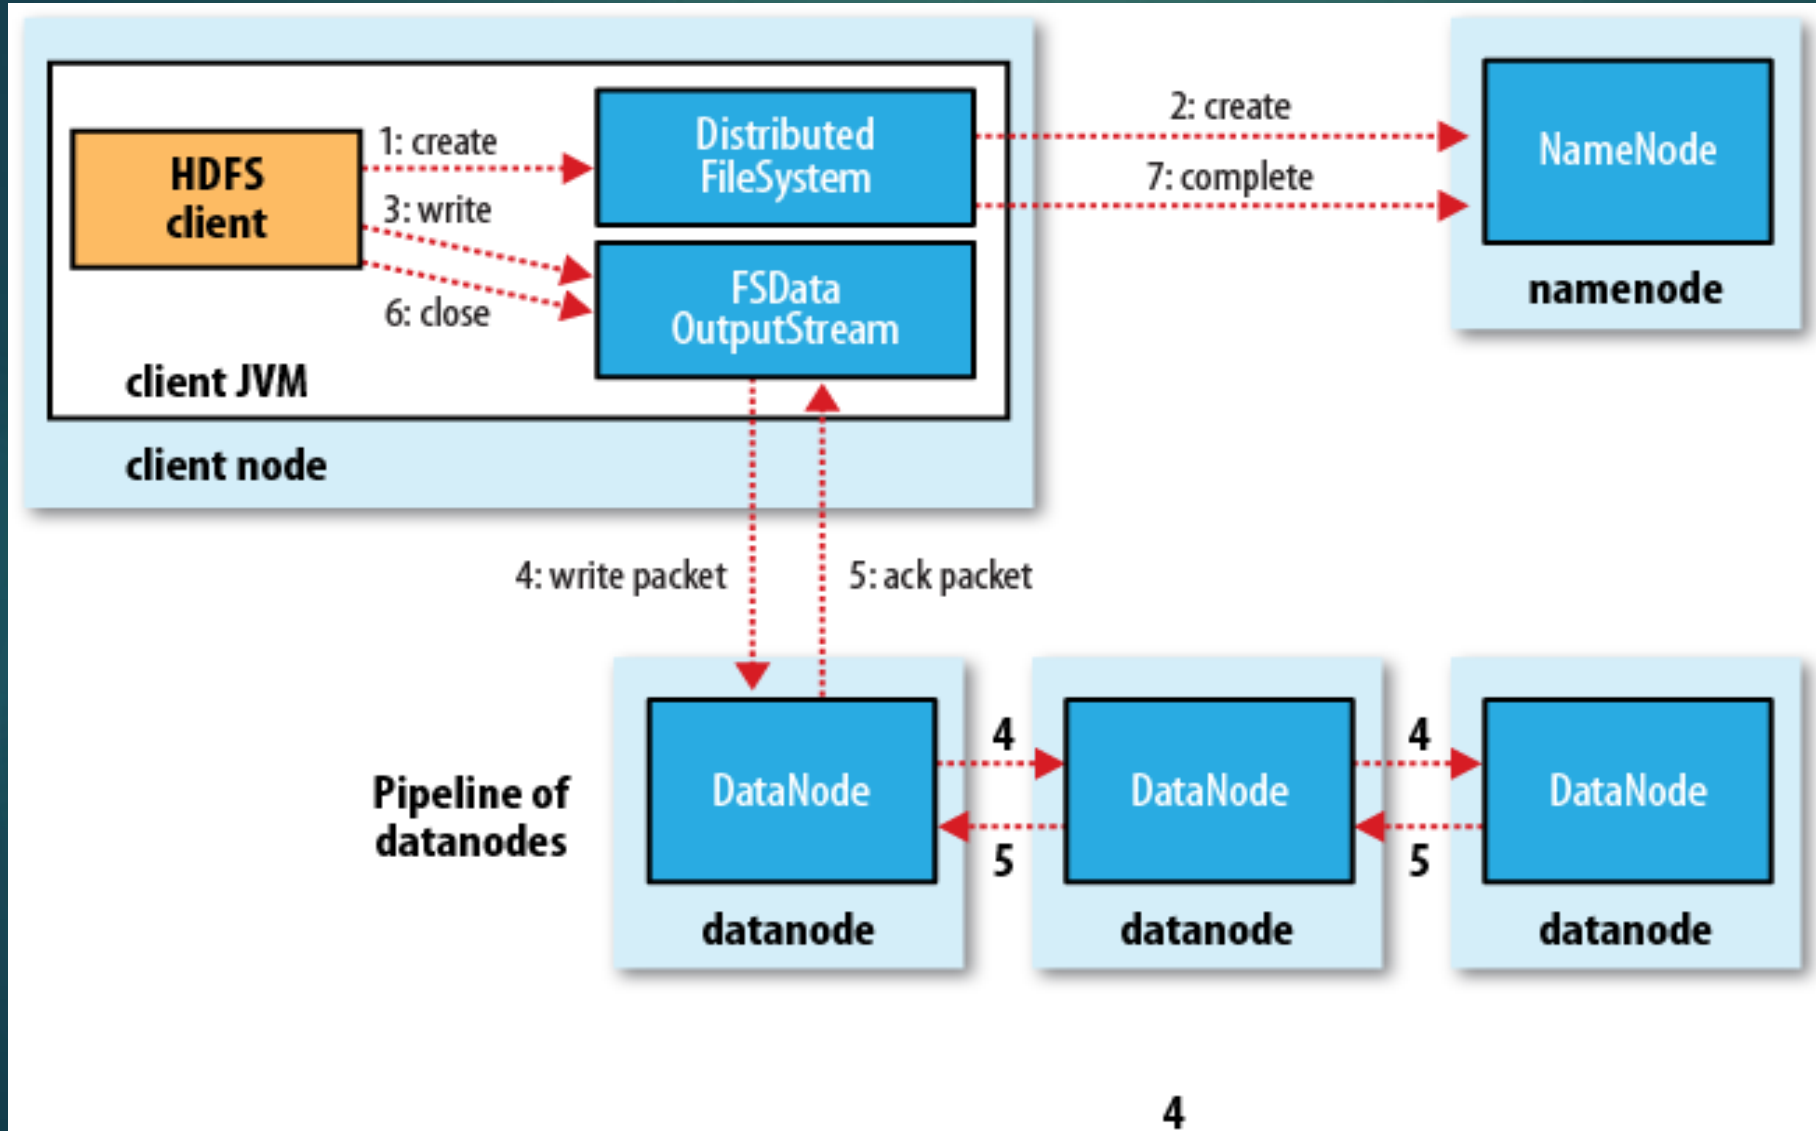
HDFS Read operation

12



HDFS Write operation

13



HDFS not made for

14

- ▶ Low-latency data access
 - $T(N) = aN + b$
- ▶ Lots of small files
 - Each file/block/directory stores around 150 bytes of metadata. Hence 1 million files each of one block will consume 300 MB of storage on Namenode
- ▶ Multiple writers, arbitrary file modifications

Exercise

15

- ▶ Calculate the memory(RAM) requirement on Namenode for given cluster configurations
 - Cluster size: 200 nodes
 - Storage capacity of each node: 24 TB
 - Block size: 128MB
 - Replication factor: 3
 - Metadata storage size for each block: 150 bytes
- ▶ Solution
 - $(200 * 24 * 10^{12} * 150) / (128 * 10^6 * 3)$

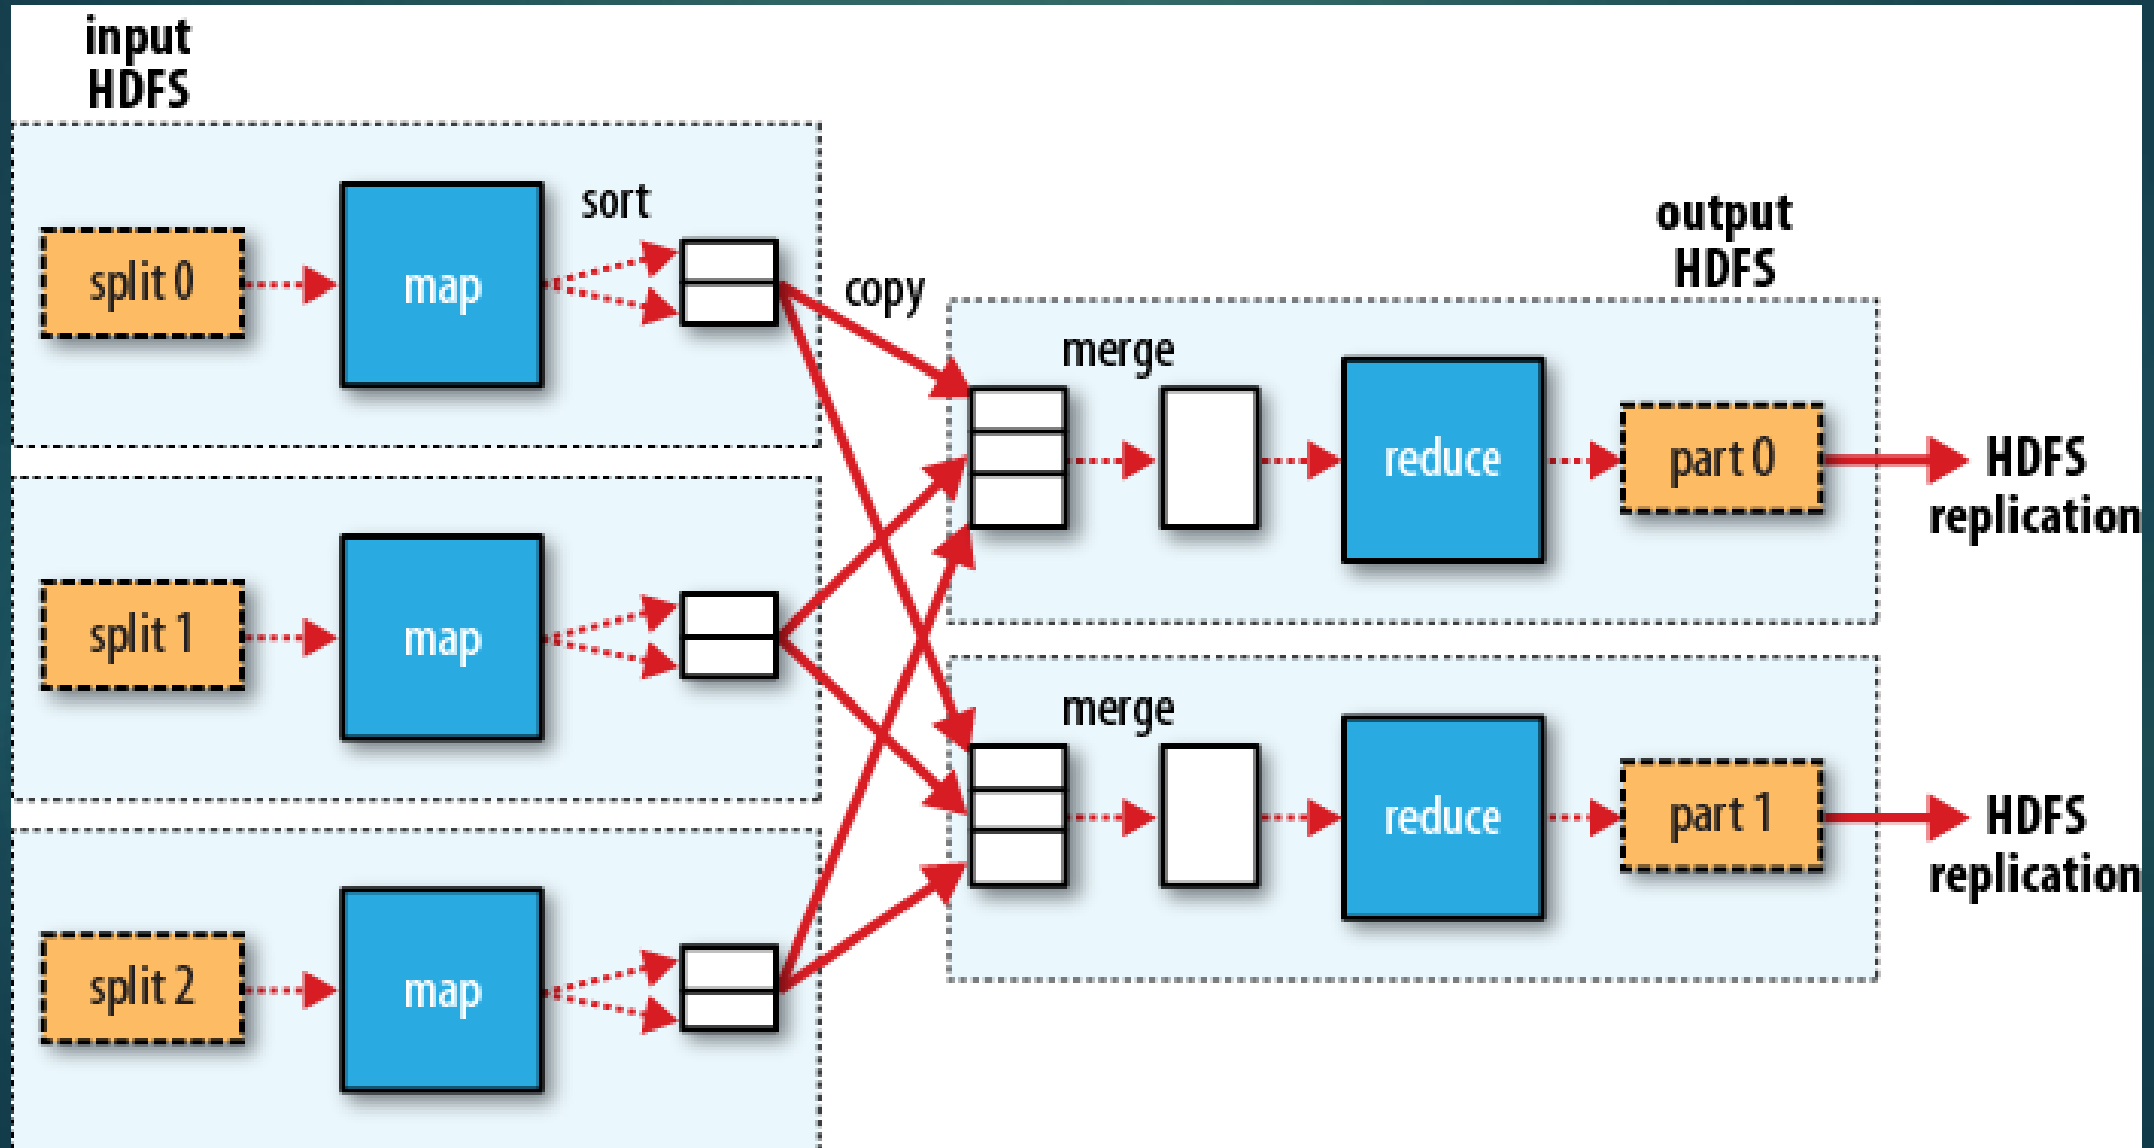
MapReduce

16

- ▶ Two major phases: Map and Reduce
- ▶ Notion of <Key, Value> pairs
- ▶ Divides job into multiple tasks
- ▶ Map: extract important information from each record
- ▶ Reduce: Aggregate, Summarize, Filter, Transform

MapReduce Stages

17




Map Tasks

18

- ▶ What is a good approach to decide how many map tasks a job should launch?
- ▶ Less number of big tasks
vs
higher number of small tasks
- ▶ Normally same as input data blocks

Task to node mapping

- Notion of data locality



Input formats


20

Input format	Description
TextInputFormat	Read Text file line by line. Key is offset and value is record text
KeyValueTextInputFormat	Tab separated key values from a text file
SequenceFileInputFormat<K,V>	Hadoop's file format

Input Splits

21

- ▶ Blocks are of fixed size
- ▶ Good chances of records being split between two block



Reduce Tasks

22

- ▶ Can be configured by programmer
- ▶ Normally same as #datanodes participating in execution
- ▶ Input Key and Value type should be same as output type of mapper
- ▶ One output file per reducer under output directory
- ▶ Generates exception if output directory already exists.
Any guess why?

Output Formats

23

Output Format	Description
TextOutputFormat<K,V>	Tab separated key value pairs in plain text format. One record per key value pair
SequenceFileOutputFormat<K,V>	Hadoop's Sequence file format
NullOutputFormat<K,V>	Nothing. Helps in map only job

MapReduce: Mapper code

```
public class WebHitCounterMapper extends
Mapper<Input Key, Input Value, Output Key, Output Value>
{

    public void map(Input Key, Input Value, Context context)
throws IOException, InterruptedException {

        <MAP Logic goes here>

        context.write(Output Key, Output Value)
    }
}
```

MapReduce: Reducer code

```
public class WebHitCounterReducer extends
Reducer<Input Key, Input Value, Output Key, Output Value>
{
    public void reduce(Input Key, Iterable<Value Data type>
values, Context context) throws IOException,
InterruptedException {

    <REDUCE logic goes here>

    context.write(Output Key, Output Value);
}
}
```

MapReduce: Driver Code

26

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
    }  
}
```

MapReduce: Driver Code

27

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
        job.setJarByClass(main.WebHitCounterMain.class);  
        job.setMapperClass mapper.WebHitCounterMapper.class);  
        job.setReducerClass(reducer.WebHitCounterReducer.class);  
  
    }  
}
```

MapReduce: Driver Code

28

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
        job.setJarByClass(main.WebHitCounterMain.class);  
        job.setMapperClass mapper.WebHitCounterMapper.class);  
        job.setReducerClass(reducer.WebHitCounterReducer.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
    }  
}
```


MapReduce: Driver Code

29

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
        job.setJarByClass(main.WebHitCounterMain.class);  
        job.setMapperClass mapper.WebHitCounterMapper.class);  
        job.setReducerClass(reducer.WebHitCounterReducer.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
    }  
}
```

MapReduce: Driver Code

30

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
        job.setJarByClass(main.WebHitCounterMain.class);  
        job.setMapperClass(mapper.WebHitCounterMapper.class);  
        job.setReducerClass(reducer.WebHitCounterReducer.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}
```

Programming Exercise

31

Challenges with Hadoop 1

32

- ▶ Applications were limited to MapReduce implementations only
- ▶ Namenode machine crash or maintenance activity
- ▶ Namespace scaling
- ▶ Backup and Recovery
- ▶ Batch oriented architecture
- ▶ Support for various file formats
- ▶ Dual responsibilities of Job tracker

Image Ref: <https://www.greycampus.com/blog/big-data/top-differences-between-hadoop-1-0-and-hadoop-2-2>

Hadoop 2

33

- ▶ Support for other data processing engines
- ▶ High Availability
- ▶ HDFS Federation
- ▶ HDFS Snapshot
- ▶ Introduced Streaming and Interactive analysis
- ▶ Support for various file formats
- ▶ Yarn

► Yet Another Resource Negotiator

MapReduce 1	YARN
Job Tracker	Resource Manager, Application Master and Timeline server
Task Tracker	Node Manager
Slot	Containers

YARN model

35

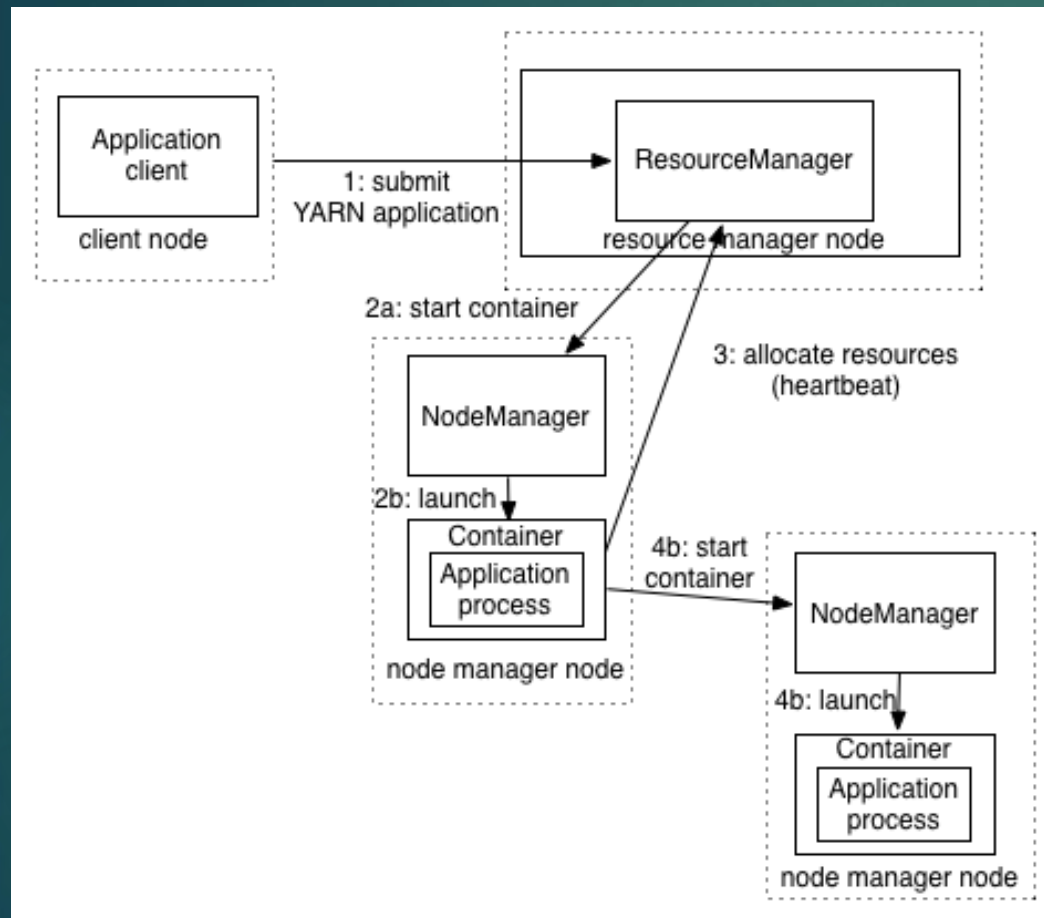


Image Ref: Hadoop definitive guide 4th edition

Pros of YARN

36

- ▶ Scalability
- ▶ Availability
- ▶ Utilization
- ▶ Multitenancy

Reference

37

- ▶ Image reference: Hadoop definitive guide 4th edition

Additional resources

38

- ▶ Reference Book
 - Hadoop: The definitive guide by Tom White ([Weblink](#))
- ▶ MapReduce google paper
 - <https://ai.google/research/pubs/pub62>
- ▶ Filesystem metadata on Namenode
 - <https://hortonworks.com/blog/hdfs-metadata-directories-explained/>
- ▶ Rack awareness
 - <https://data-flair.training/blogs/rack-awareness-hadoop-hdfs/>
 - <https://community.hortonworks.com/articles/43057/rack-awareness-1.html>
- ▶ Concurrency vs Parallelism
 - <https://stackoverflow.com/questions/1050222/what-is-the-difference-between-concurrency-and-parallelism>
- ▶ Latency vs Throughput
 - <https://stackoverflow.com/questions/16718095/high-throughput-vs-low-latency-in-hdfs>